

- Last time: “best effort” delivery as the service abstraction
  - Not delivered -> [timeout + retransmit](#)
  - Delivered n > 1 times -> [transform operations to be idempotent](#)
  - Delivered altered -> **checksum or crypto**
  - Delivered out of order -> [sequence number](#)
  - On top of this service abstraction, we can build:
    - VoIP
    - User Datagrams
    - VPN (IP-in-UDP/IP-in-IP/IPsec)
      - Q: How does Netflix determine where an IP address is actually from?
      - A: Netflix would look at the IP addresses provided by VPN services and ban those IP addresses.
- Short get: get(key) -> value
  - E.g. host: what is the IP address that corresponds to a host?
    - With packet loss, it takes a longer time to reply, but would still give an answer
    - This service is “reliable” despite the fact that it is built on a unreliable “best effort” service abstraction

C/C++

```
// Server
void recv ( const string& service ) {
    UDPSocket sock;
    sock.bind ( Address ( "0", service ) );
    Address source ( "0" );
    string payload;
    while (true) {
        sock.recv( source, payload);
        cout << "Message from" << source.to_string() << ": "
        << payload << endl;
        if (payload == "best_class_ever" ) {
            sock.sendto( source, "EE180");
        }
    }
}
```

C/C++

// Sender

```
void run( const string& host, const string& service, const
string& query) {
    UDPSocket sock;
    sock.set_blocking( false );
    Address source ( "0" );
    string answer;

    // retransmit the query (with a small timeout), until there
is a reply
    do {
        sock.sendto(Address(host, service), query);
        this_thread::sleep_for(seconds(1));
        sock.recv(source, answer);
        if (answer.empty()) {
            cerr << "No reply, retransmitting" << endl;
        }
    } while (answer.empty())

    cout << "Got reply to " << query << ": " << answer << endl;
}
```

- By doing this, we implement a “reliable” service on top of an “unreliable” service abstraction, and this is also how many real-world reliable services are built (e.g. host).
  - And also: Domain Name System (DNS): what is the IP address of an internet domain name?
  - DHCP (Dynamic Host Configuration Protocol): what is the IP address I am supposed to use?
- Set: (e.g. set the back door open)
  - Both short get and set (the back door open), you could say how ever many times you want and it does not change the ending state
  - But for ``pop(7)``, ``push("hi")``, it matters how many times you say it.
  - Idempotent: doing one time or more than one time does not change the ending state (GET PUT). The strategy we used above works for something idempotent, but not for non-idempotent action
- Do a non-idempotent operation (POST):
  - By having a set of launched missiles, we make launch\_missile idempotent

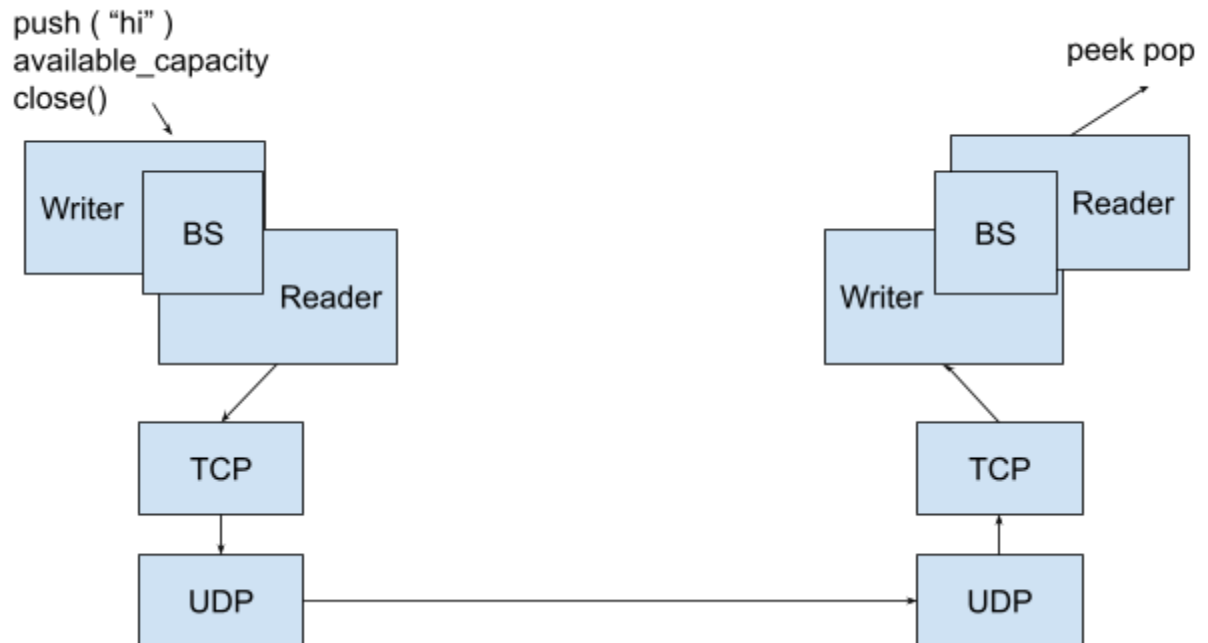
C/C++

// Server

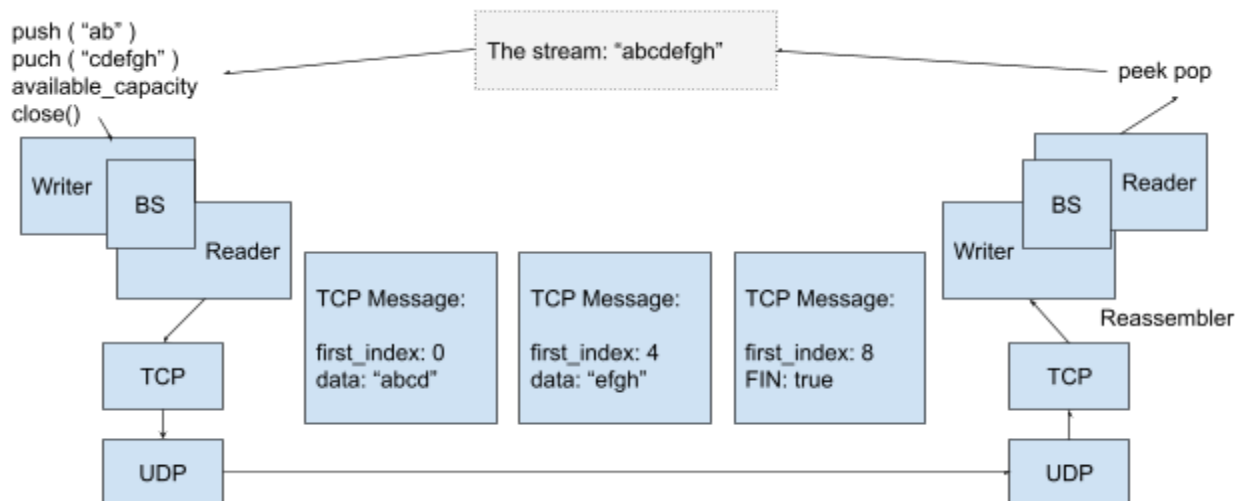
```
void launch_missile() {
    cout << "Launching one missile" << endl;
}

void recv ( const string& service ) {
    unordered_set<uint64_t> launched_missile;
    UDPSocket sock;
    sock.bind ( Address ("0", service) );
    Address source ("0");
    string payload;
    while (true) {
        sock.recv( source, payload);
        cout << "Message from" << source.to_string() << ": "
        << payload << endl;
        if (payload == "best_class_ever" ) {
            sock.sendto(source, "EE180");
        } else if (payload == "launch_one_missile" + missile_id
        ) {
            if (missile_id not in launched_missile ) {
                launch_missile();
                launched_missile.insert(missile_id);
            }
            sock.sendto(source, "ack");
        }
    }
}
```

- ByteStream: push, pop, peek needs to be transformed into idempotent operations, and this is achieved by **TCP**



- 
- What should be in the TCP Sender message to make these operations idempotent?
  - `push("abcd")` works iff each message is delivered exactly once
  - `push("abcd") + message unique id`, but the sender needs to keep a set of any message sent
  - Create a reassembler, `first\_index: 0, data: "abcd"` `first\_index: 4, data: "efgh",`  
`first\_index = 8, FIN=true`

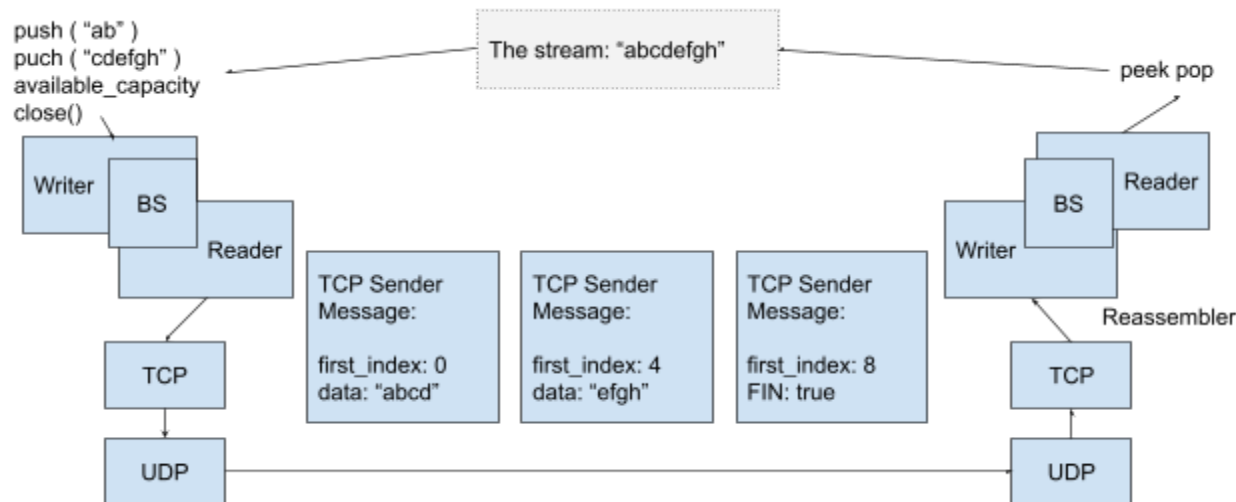


- Stacks of service abstraction
  - Short gets (DNS, DHCP) -> User datagrams -> Internet Datagrams
  - Byte Stream -(TCP)--> User datagrams -> Internet Datagrams
    - Web requests/responses (HTTP) -> Byte Stream
      - Youtube/Wikipedia -> Web requests/responses
    - Email sending (SMTP) -> Byte Stream

- Email receiving (IMAP) -> Byte Stream
- Multiplexing ByteStream
  - “u8 u8” (Which byte stream; what is the byte)
    - Any reading and writing of one byte would be actually two bytes, the first byte for which byte stream and the second for the actual byte
  - “u8 u8 {u8 u8 ... u8}” (which byte stream; size of payload; sequence of characters of the string chunk)
    - Any reading and writing of n bytes would be actually n + 2 bytes
    - Tagged byte stream: HTTP/2 | SPDY
- How to make ByteStream push idempotent?

- TCP Sender Message

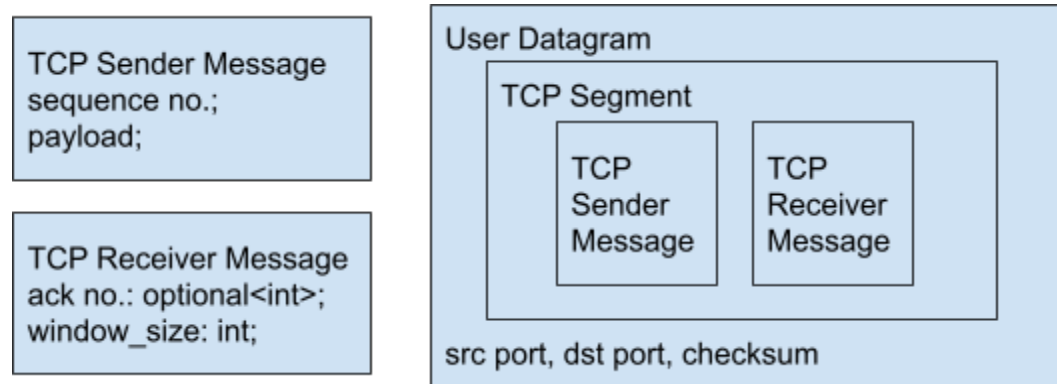
first\_index  
data  
FIN



- This works for out-of-order or multiple deliveries. Since UDP has a checksum field, altered TCP Message would be ignored on the UDP layer.
- What if datagrams are missing?
  - How does the sender know that a datagram needs to be sent multiple times?
  - DNS/DHCP: if we don't receive an answer, then we retransmit. But such response/answer does not exist in `pushing` (`void push()`)
- Acknowledgement
  - TCP Receiver Message
    - “A B C D E F G” each byte sent as a separate TCP sender message, and “D” is not received.
    - “I got the sender message with first-index = 2, length = 1.”
      - Valid but more work. There will be one receiver message for each sender message.
    - “Got anything? Y/N. Next needed: #3.”
      - Acknowledgements are accumulative, and that make life simpler.

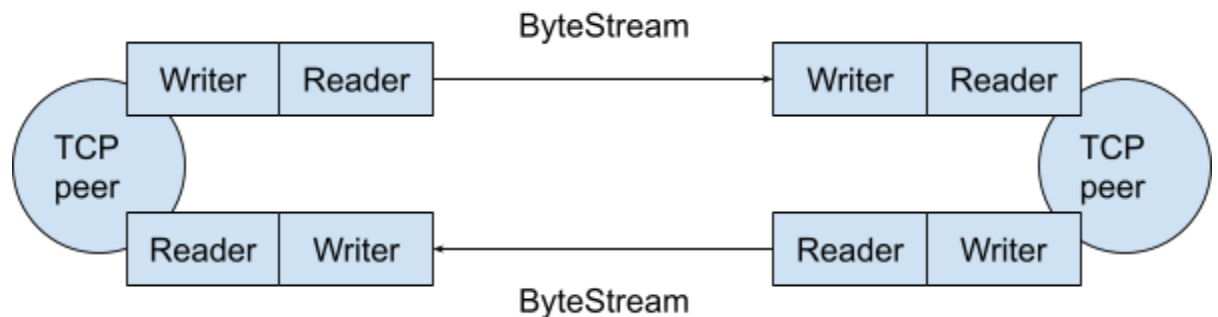
- Give FIN flag a number: "A B C D E F G FIN".  
TCP Sender Message: {sequence number, data}  
TCP Receiver Message: {Next needed: optional<int>}  
and TCP Receiver Message {Next needed: optional(8)} would mean FIN is received.
- TCP Receiver Message: {Next needed: optional<int>; available capacity:int}
  - {Next needed: optional(3); available capacity: 2} === Receiver wants to here about "DE".
  - "DE" is the **window**. (Red area in that picture of lab 1)
- **TCP Receiver Message: {Ack no: optional<int>; window size: int}**

- TCP Segment

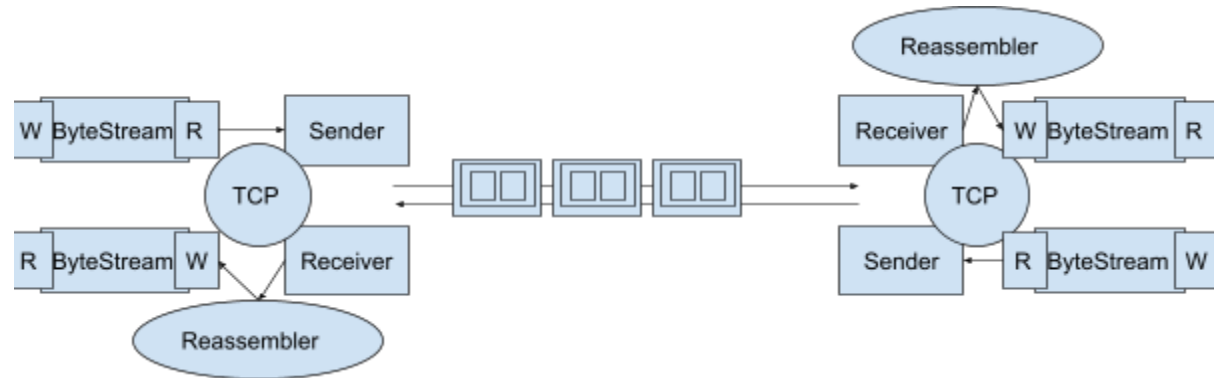


- This is the service abstraction that TCP is providing:

#### Service abstraction of ByteStream

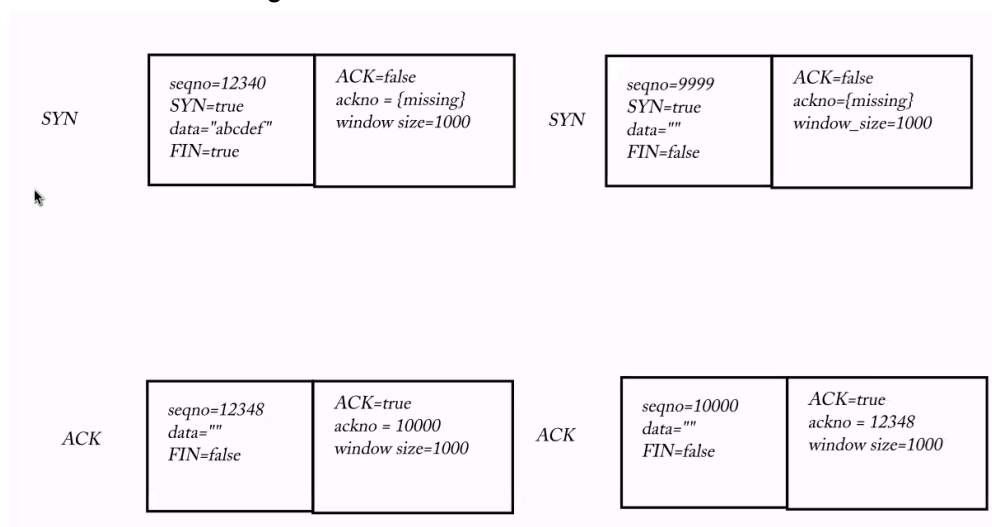


- And this is what happens under the hood (and also what you will be implementing in the labs)



- Rules of TCP
  - Reply to any nonempty TCP Server Message
- What happens when a stream ends?
  - My sender has ended its outgoing bytestream, but the incoming bytestream from the peer may not be ended.
  - When a stream ends, can the same pair of ports be used? Reusing the same pair of ports makes it not clear to tell whether a datagram belongs to the old stream or the new stream.
  - We want a new INCARNATION of the connection (new connection on the same pair of ports)
  - **Sequence number**: start from a random big number + **SYN**: this sequence number should be viewed as the beginning of a stream
    - If the sequence number doesn't make sense on the old stream, and the SYN flag is true, the receiver knows this is a new incarnation of the connection.
    - e.g. {sequence\_no=12345, data="abcdef", SYN=true, FIN=true}, and {sequence\_no=99999, data="xyz", SYN=true, FIN=true}
  - First seqno belongs to SYN flag, next seqnos belong to each byte of stream, final seqno belongs to FIN flag.
    - It is very important to have SYN flag and FIN flag delivered reliably, so therefore receiver need to acknowledge SYN seqno and FIN seqno
- What happens to TCP receiver message's next\_needed\_idx field before receiving the SYN flag from the peer?
  - Without seqno:
    - I: {{first\_index=0, data="abcdef", FIN=true}, {next\_needed=0, window\_size=1000}}
    - Peer: {{first\_index=0, data="", FIN=true}, {next\_needed=7, window\_size=1000}}
    - I: {{first\_index=7, data="", FIN=false}, {next\_needed=1, window\_size=1000}}
  - With seqno and SYN:

- I: {{seqno=12340, SYN=true, data="abcdef", FIN=true}, {**What should this be? (before seeing 9999 from the Peer)**}}
- Peer: {{seqno=9999, SYN=true, data="", FIN=true}, {next\_needed=12348, window\_size=1000}}
- I: {{next\_needed=10001, window size =1000}}
- ackno = optional<int> (a pair of ACK flag and ackno int)
  - I: {{seqno=12340, SYN=true, data="abcdef", FIN=true}, {ACK=false, ackno={missing}, window\_size=1000}} (SYN)
  - Peer: {{seqno=9999, SYN=true, data="", FIN=true}, {ACK=true, ackno=12348, window\_size=1000}} (SYN+ACK)
  - I: {{ACK=true, ackno=10001, window size =1000}} (ACK)
- (SYN) + (SYN+ACK) + (ACK) = "the three-way handshake"
- What if the two SYN messages are sent at the same time?



- Not a classic "three-way handshake" but still a valid way of starting a TCP connection.
- Standardized TCP Message:
  - **Sender:** {sequence number, SYN, data, FIN}
  - **Receiver:** {ackno: optional<int>, window\_size}
  - "User Datagram" info

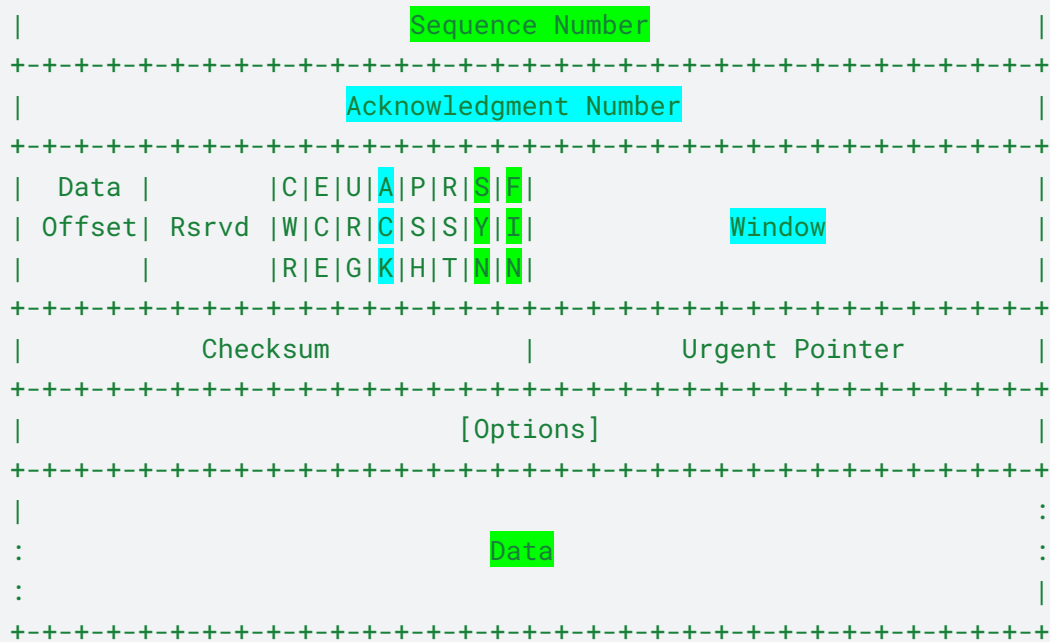
<https://www.rfc-editor.org/rfc/rfc9293.html#name-header-format>

Unset

```

0          1          2          3
0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5 6 7 8 9 0 1
+-----+-----+-----+-----+
|                Source Port          |                Destination Port          |
+-----+-----+-----+-----+
```





Note that one tick mark represents one bit position.

- Wireshark tool